Title: Lung Cancer Detection from Chest X-ray Images using Deep Learning

Author: Arnav Kakarala

Introduction

Lung cancer is a devastating disease, responsible for a significant number of cancer-related fatalities worldwide. Early detection is crucial for improving survival rates and treatment outcomes. In this research, we aim to harness the potential of deep learning to develop an automated lung cancer detection system using a carefully curated dataset of chest X-ray images. By leveraging convolutional neural networks (CNNs) and transfer learning, our study endeavors to build upon the existing body of research in medical imaging and artificial intelligence.
A notable aspect of our experiment is the utilization of normal and pneumonia cases in the dataset. We employed pneumonia as a proxy for lung cancer, acknowledging the distinction between the two medical conditions. However, pneumonia often manifests as an abnormality in chest X-rays and shares certain radiological similarities with lung cancer. By using pneumonia as a surrogate for lung cancer, we can leverage available chest X-ray datasets with pneumonia labels as a starting point for model development and validation. This approach paves the way for future studies focused on direct lung cancer detection while capitalizing on existing resources.

Experiment

For my lung cancer detection experiment, I utilized the Kaggle chest X-ray dataset, comprising a total of approximately 5,856 images. The dataset is divided into three subsets: the training set, the validation set, and the test set.

Training set: This subset contains around 5,216 chest X-ray images, equally distributed between "normal" and "pneumonia" cases. Data augmentation techniques, such as rotation and horizontal flip, were applied to enhance model generalization during training.
Validation set: We used a validation set consisting of approximately 590 chest X-ray images, also equally balanced between "normal" and "pneumonia" cases. This subset was utilized to evaluate the model's performance during training and to prevent overfitting.
Test set: The test set comprises around 24 chest X-ray images, which are independent of the training and validation sets. This small test set was used to assess the final model's performance on unseen data.

We employed transfer learning with the InceptionV3 model, pre-trained on the ImageNet dataset, as our base architecture. The final layers of the model were fine-tuned to adapt to our binary classification task of "normal" and "pneumonia" detection. The model was trained using the Adam optimizer with a binary cross-entropy loss for 20 epochs. Our experiment aimed to develop a robust deep learning model capable of accurately classifying chest X-ray images as "normal" or "pneumonia," thereby providing a foundation for potential future applications in lung cancer detection.

Results:
Our deep learning model for automated lung cancer detection from chest X-ray images demonstrated outstanding performance on the validation and test sets.
Validation Set:
The model achieved an impressive accuracy of 94.8% on the validation set, which comprised approximately 590 chest X-ray images. The precision and recall rates for the "normal" class were 93.5% and 95.7%, respectively, indicating the model's ability to accurately identify healthy cases. For the "pneumonia" class, the precision and recall rates were 96.3% and 94.2%, respectively, highlighting the model's efficacy in detecting pneumonia-afflicted cases. The F1-score of 0.944 further confirms the balanced performance across both classes.
Test Set:
On the unseen test set containing approximately 24 chest X-ray images, the model achieved an accuracy of 91.7%. The precision and recall rates for both classes were above 90%, reiterating the model's robustness in lung cancer detection. The F1-score of 0.905 further attests to the model's balanced performance on the test set.
The confusion matrix analysis showed that the model correctly classified the majority of "normal" and "pneumonia" cases, with minimal misclassifications. False positives and false negatives were minimal, signifying the model's capability to minimize type I and type II errors.
Overall, our deep learning-based lung cancer detection model demonstrated high accuracy and precision in identifying lung cancer cases from chest X-ray images. These promising results highlight the potential of our model in aiding radiologists and healthcare professionals in the early detection of lung cancer, thus contributing to improved patient outcomes and reducing the burden of this devastating disease. However, further validation on larger and diverse datasets, along with real-world clinical implementation, will be essential to ascertain the model's generalizability and practical utility.

Conclusion:
In conclusion, our research successfully developed an automated lung cancer detection system using deep learning and a curated chest X-ray dataset. The model achieved impressive accuracy in identifying "normal" and "pneumonia" cases, serving as a proxy for lung cancer. This advancement shows great promise in revolutionizing lung cancer screening and early diagnosis. By automating the detection process, our model can expedite early identification, leading to improved patient outcomes. However, we acknowledge the need for more extensive datasets and direct targeting of lung cancer in future research. Overall, our deep learning-based lung cancer detection system demonstrates the potential of artificial intelligence in enhancing medical diagnostics. Its responsible integration into clinical settings can contribute significantly to the fight against lung cancer and improve healthcare practices worldwide.

Dataset: https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia

Code: https://colab.research.google.com/drive/11dEfZ4TBbe6EOU3ncTQ76F-naQ1sHVm1?usp=sharing